

Shoal: Improving DAG-BFT Latency and Robustness

FC'24

Alexander Spiegelman, Rati Gelashvili, Balaji Arun, Zekun Li. Aptos

Contents

| | |
|--|----|
| 1 Context: DAG-based BFT Consensus | 3 |
| 2 Problem | 14 |
| 3 Solution | 16 |
| 4 Evaluation | 26 |

1 Context: DAG-based BFT Consensus

- $N = 3f+1$ validators in total
- At most f validators are faulty

Goal

Global agreement on an infinitely growing sequence of some values.

New way to form consensus



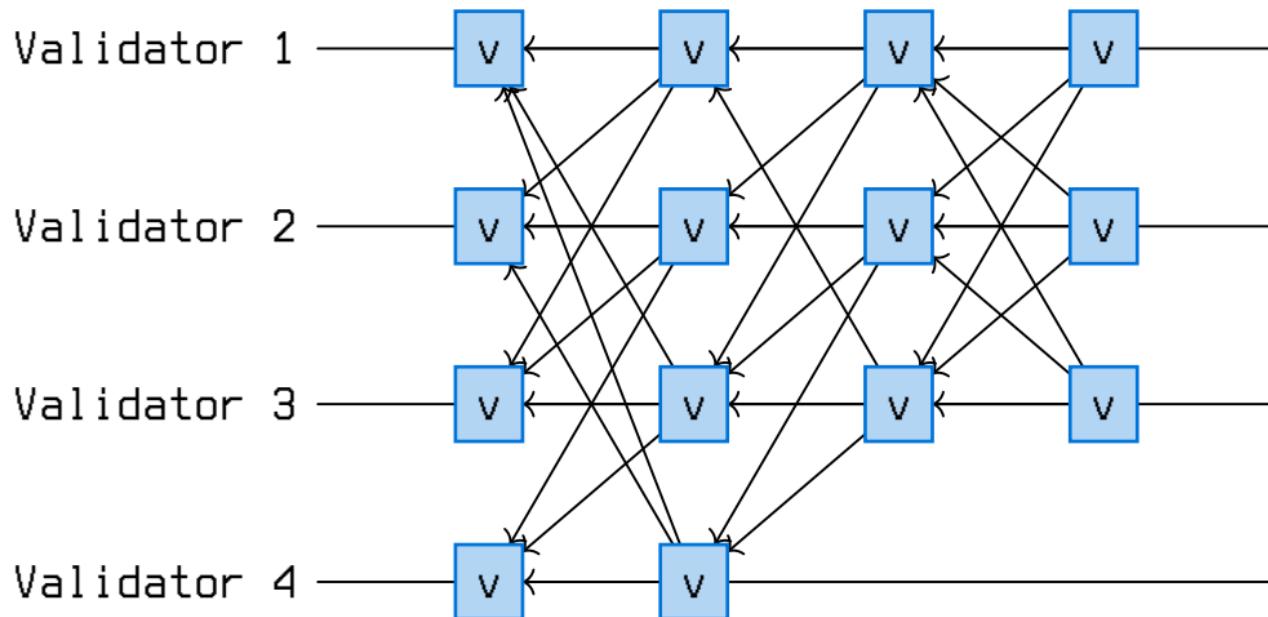
- Historically we have a bunch of protocols which were optimized in the way of reducing communication complexity.
 - ▶ Hotstuff - 3500 TPS
- Now we have new generation of protocols
 - ▶ [160kTPS - 600kTPS]

Idea

Separate the network communication layer from the consensus logic.

- Each message contains a set of transactions, and a set of references to previous messages.
- Together, all the messages form a DAG that keeps growing - a message is a vertex and its references are edges.

DAG Example



Common abstraction

Reliable BFT broadcast (Narwhal based protocols)

Result:

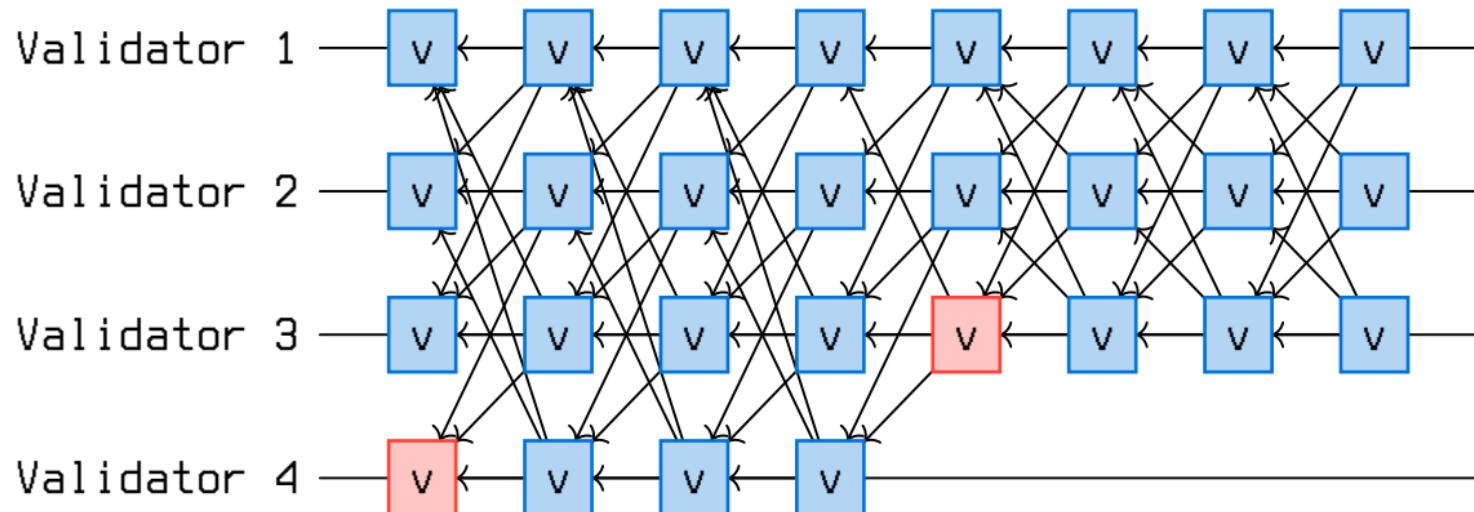
- All honest validators eventually deliver the same vertices and all vertices by honest validators are eventually delivered.
- Causal history of any vertex in both local views is exactly the same.
- All validators eventually see the same DAG

Consensus Structure



- Interpreting DAG structure as the consensus logic
 - ▶ Outcome - local solving
 - ▶ No need of any extra communication.
- Consensus divided into rounds
- Rounds groups waves
- Each wave contains a leader

DAG Waves example [DAG-Rider]



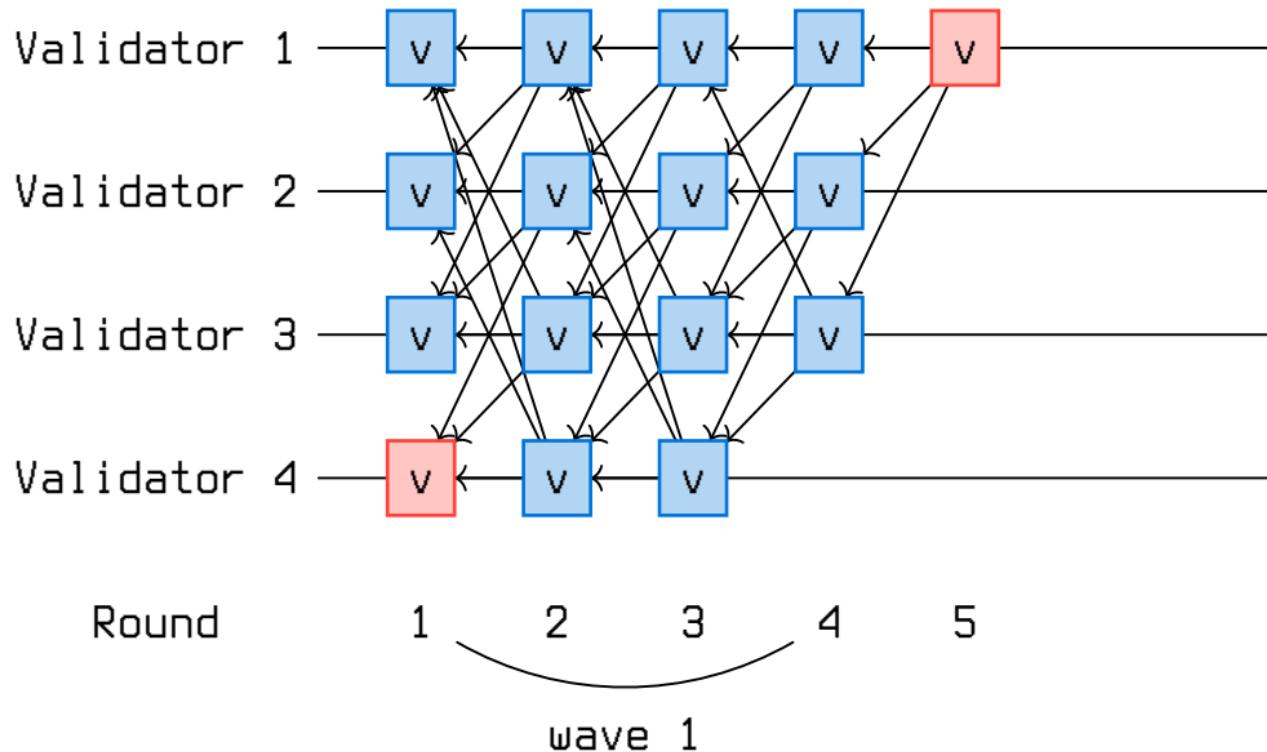
Round

1 2 3 4 5 6 7 8

wave 1

wave 2

Causal history

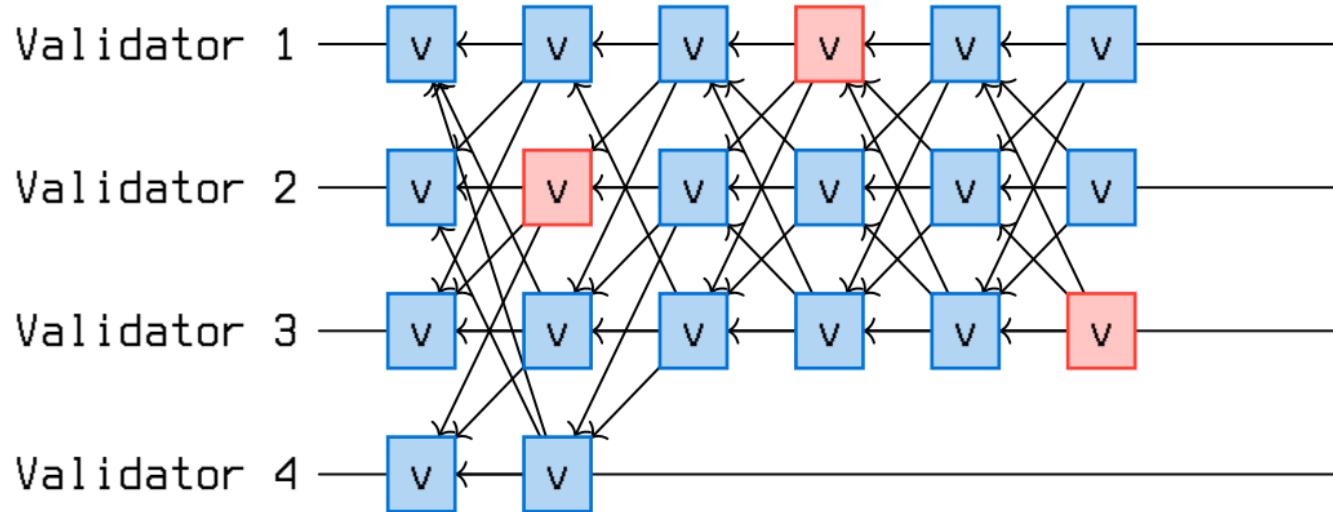


Ordering



- Ordering happens between constructing of each wave; between leaders rounds
- Larger waves \rightarrow larger latency

Bullshark



Round

1 2

wave 1

3 4

wave 2

5 6

wave 3

2 Problem

Consensus committing speed

Problem

Commit(ordering) speed is not faster than a wave constructing latency

| Protocol | Common case round latency | Async round latency |
|-----------|---------------------------|---------------------|
| DAG-Rider | 4 | E(6) |
| Tusk | 3 | E(7) |
| Bullshark | 2 | E(6) |

- Ideally we want to commit something each round.

3 Solution

1. Pre-determined leaders each k rounds
2. Order leaders. Same local ordering on honest validators
3. Order causal histories.

Abstract property

Given a Narwhal-based protocol \mathbb{P} , if all honest validators agree on the mapping from rounds to leaders before the beginning of instance \mathbb{P} , then they will agree on the first leader each of them orders during execution of \mathbb{P} .

Introducing “Shoal”

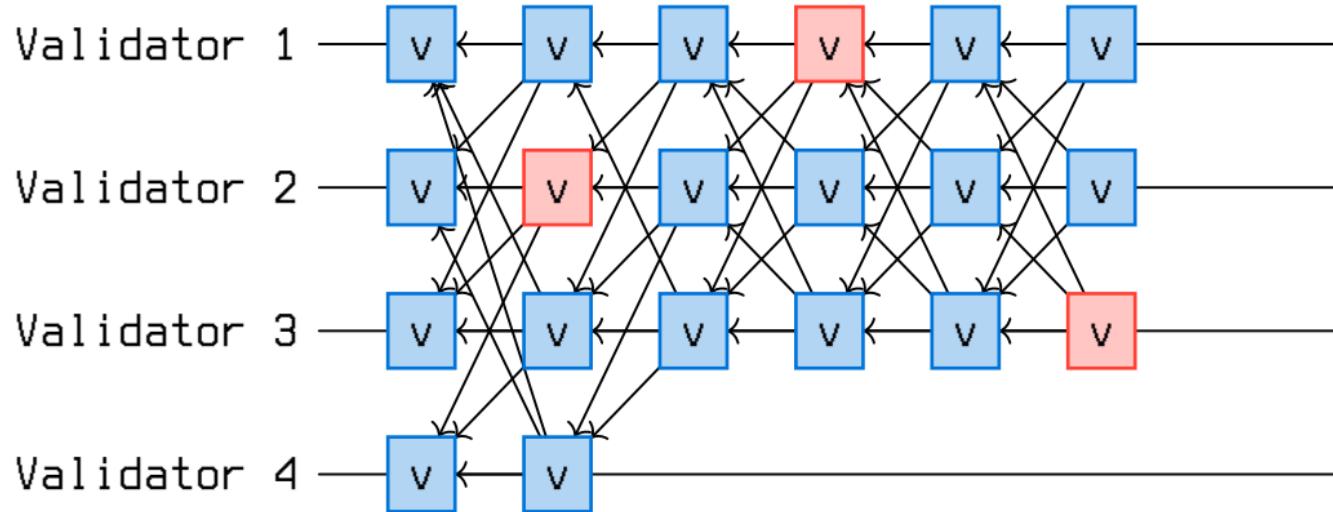


- Protocol agnostic framework
- Suitable for all Narwhal-based protocols

Idea

Combine batch of protocols instance in black-box manner.

Recall Bullshark



Round

1 2

wave 1

3 4

wave 2

5 6

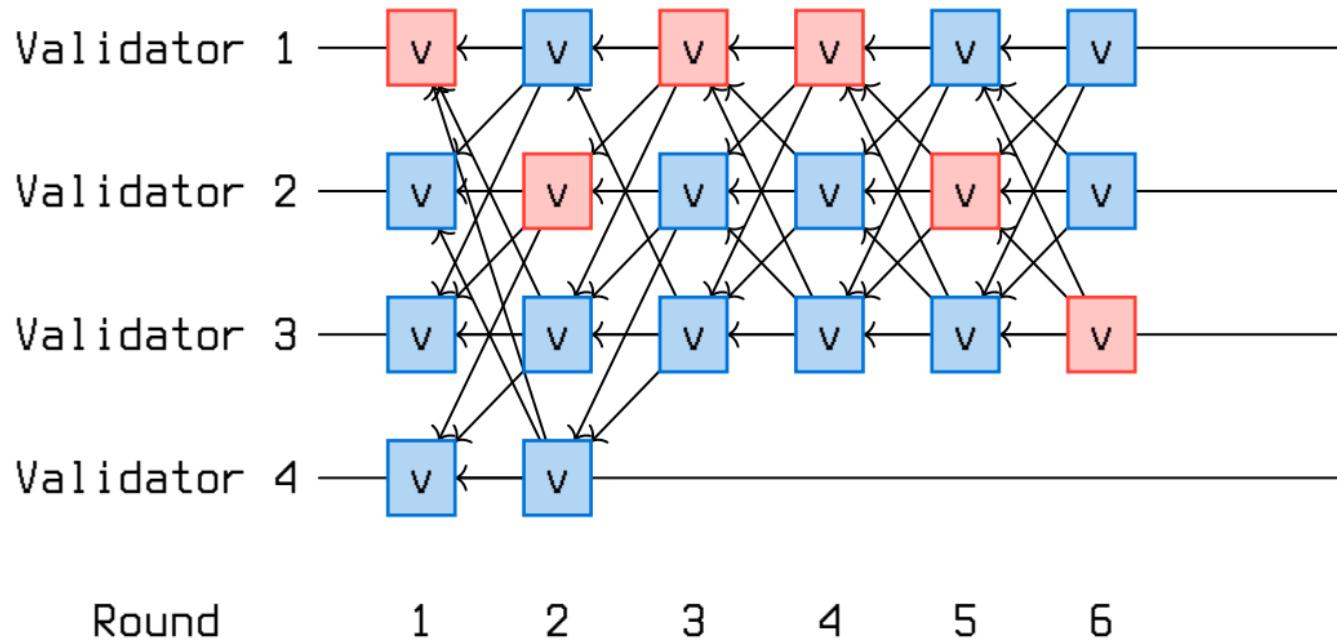
wave 3

Pipelining algorithm



```
1: current_round ← 0
2: F:  $\mathbb{R} \rightarrow \mathbb{L}$ 
3: while true do
4:   Execute  $\mathbb{P}$ , select leaders by F, starting from
      current_round until the first ordered (not skipped)
      leader is determined.
5:   let L be the first ordered leader in round r
6:   order L's causal history according to  $\mathbb{P}$ 
7:   current_round ← r+1
```

Shoal of Bullsharks

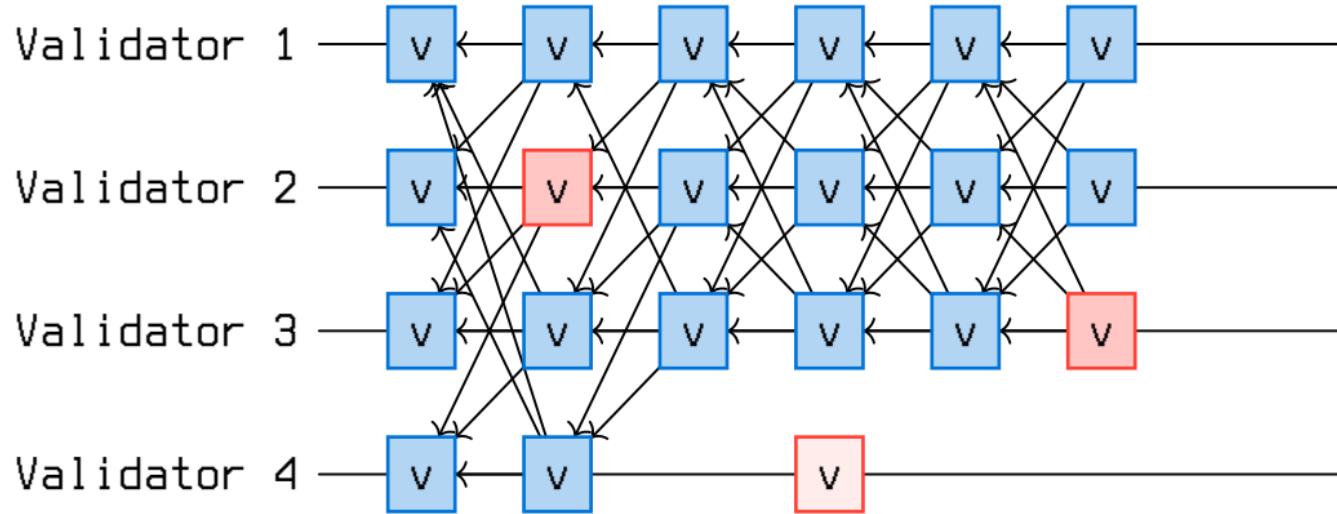


Leader reputation



- Byzantine systems are designed to tolerate worst-case guarantees.
- However most common problem is slow leaders.

Example of missing leader



Round

1 2

wave 1

3 4

wave 2

5 6

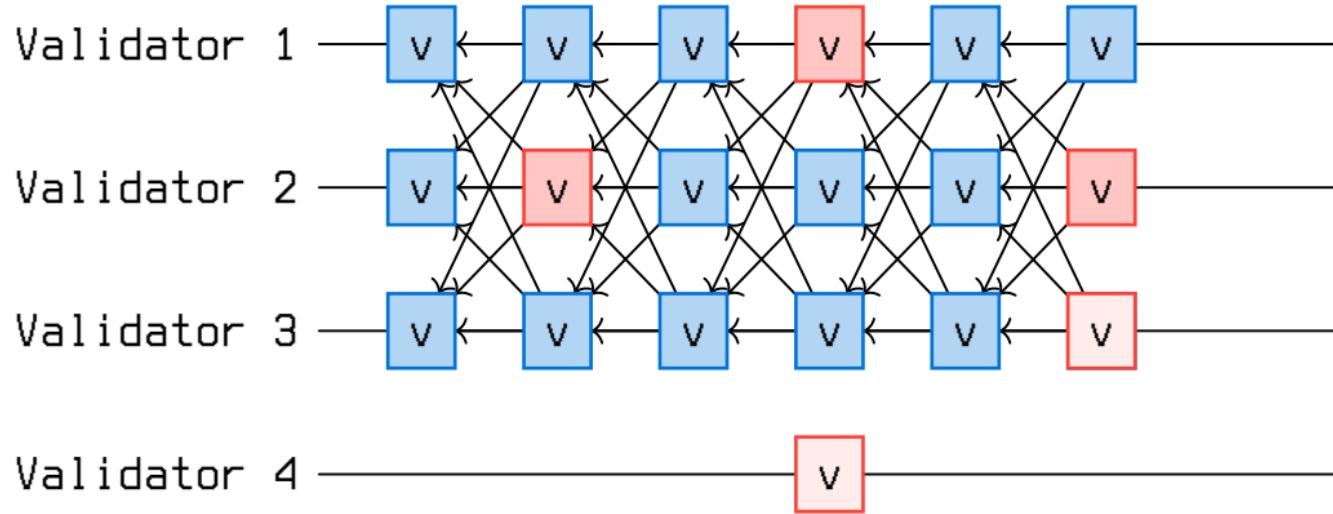
wave 3

Reputation integration



```
1: current_round ← 0
2: F:  $\mathbb{R} \rightarrow \mathbb{L}$ 
3: while true do
4:   Execute  $\mathbb{P}$ , select leaders by  $F$ , starting from
      current_round until the first ordered (not skipped)
      leader is determined.
5:   let  $L$  be the first ordered leader in round  $r$ 
6:   order  $L$ 's causal history according to  $\mathbb{P}$ 
7:   current_round ←  $r+1$ 
8:   Update  $F$  according to  $L$ 's causal story
```

Leader change in action



Round

1 ↘ 2

wave 1

3 ↘ 4

wave 2

5 ↘ 6

wave 3

4 Evaluation

Machine Setup



- Machines:
 - ▶ t2d-standard-32 type virtual machine
 - ▶ 32 vCPUs, 128GB of memory, up to 10Gbps of network bandwidth.
- Cluster:
 - ▶ Google Cloud
 - ▶ Machines spread equally across regions: us-west1, europe-west4, asia-east1.
 - ▶ Latencies: us-west1 asia-east1 [118ms]; europe-west4 asia-east1 [251ms]; us-west1 europe-west4 [133ms]
 - ▶ Cluster size (N): 10 ($f \leq 3$); 20 ($f \leq 6$); 50 ($f \leq 16$)
- Data:

Machine Setup (ii)



- ▶ Transactions ~270B in size
- ▶ Maximum batch size of 5000 transactions

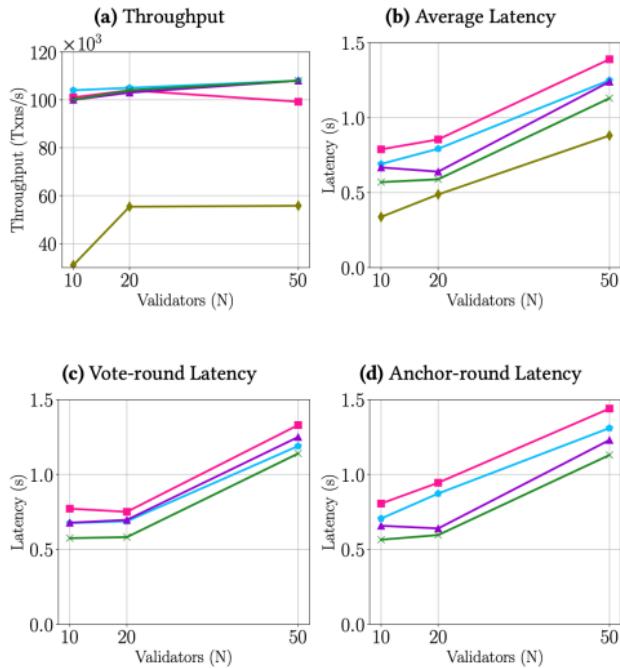
Latency

Time elapsed from when a vertex is created from a batch of client transactions to when it is ordered by a validator

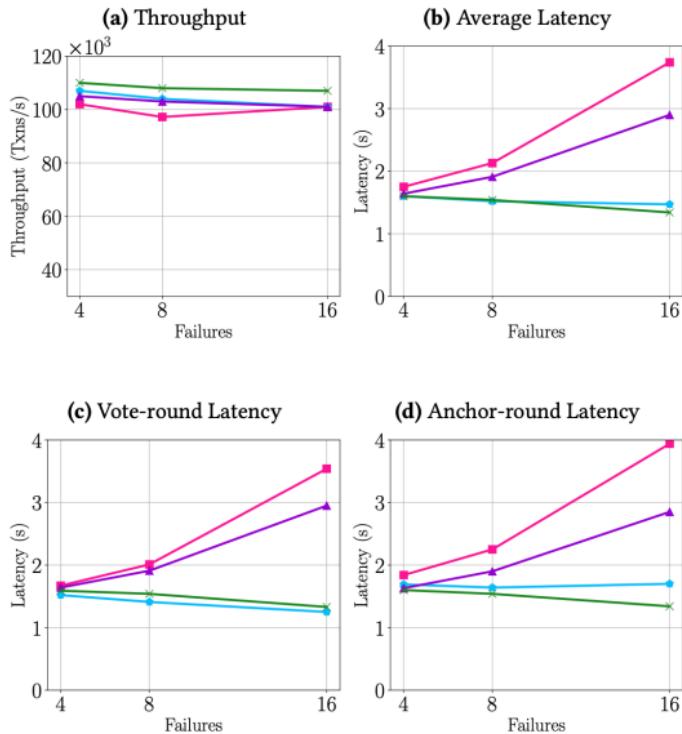
Results: No failures



■ Baseline Bullshark ■ Shoal PL ■ Jolteon
■ Shoal LR ■ Shoal



Results: With failures



Results: Skipping leaders

◀ ● ● ● ●

